

# CS550 “Advanced Operating Systems”

## Instructor: **Professor Xian-He Sun**

- Email: [sun@iit.edu](mailto:sun@iit.edu)
  - Office: SB235C
  - Class time: Monday, Wed., 3:15pm-4:30pm, **HH MEZZANINE**
  - Office hour: Monday, Wednesday, 4:45-5:45pm
  - <http://www.cs.iit.edu/~sun/cs550.html>
- 
- TA: Mr. Hua Xu, Email: [hxu40@hawk.iit.edu](mailto:hxu40@hawk.iit.edu)
  - Office Hour: 11am - 12pm, Tuesday
  - [meet.google.com/kfp-pysg-cat](https://meet.google.com/kfp-pysg-cat)
  - Office Hour: 12pm - 1pm, Tuesday & Thursday  
[meet.google.com/bnn-eqao-htg](https://meet.google.com/bnn-eqao-htg)
- 
- Blackboard:
    - <http://blackboard.iit.edu>
  - Substitute lecturer:
    - Anthony Kougkas, assistant research professor
    - [akougkas@hawk.iit.edu](mailto:akougkas@hawk.iit.edu)

# Chapter 3: Processes

- Overview of processes and threads
- Virtualization
- Clients
- Servers
- Code/Process migration

# Client & Server

# Outline

- Networked user interfaces
- Thin-client network computing
- Client-side software for distribution transparency
- Concurrent versus iterative servers
- End-points
- Interrupting a server
- Stateless vs stateful servers
- Server clusters

# Concurrent vs Iterative servers



- **Iterative server:**
  - the server itself handles the request and, if necessary, returns a response to the requesting client
- **Concurrent server:**
  - does not handle the request itself but passes it to a separate thread or another process, after which it immediately waits for the next incoming request.
  - Multithreaded: fork a new process for each new incoming request

# Contacting a server

- Clients send requests to an **end point**, also called a **port**, at the machine where the server is running.
- Each server listens to a specific end point.
- How do clients know the end point of a service?
  - Globally assign end points for well-known services
    - FTP port 21, HTTP port 80 etc
  - client needs to find only the network address of the machine where the server is running → Name services

# Contacting a server

- Dynamically assigned to a service by its local OS
  - special daemon running on each machine that runs servers
  - daemon keeps track of the current end point of each service implemented by a co-located server
  - The daemon itself listens to a well-known end point.
  - Use **superservers** for better resource utilization



# Interrupting a server

- Whether and how a server can be interrupted?
  - abruptly exit the client application (which will automatically break the connection to the server),
  - immediately restart it and pretend nothing happened.
  - The server will eventually tear down the old connection, thinking the client has probably crashed.
- Better solution:
  - **Out-of-band:** data that is to be processed by the server before any other data from that client.
    - server listens to a separate control end point to which the client sends out-of-band data, while at the same time listening (with a lower priority) to the end point through which the normal data passes.
    - send out-of-band data across the same connection through which the client is sending the original request → Urgent data



# Stateless vs Stateful servers

- A **stateless server** does not keep information on the state of its clients, and can change its own state without having to inform any client.
  - e.g. Web server:
    - responds to incoming HTTP requests (store or fetch some data)
    - When the request has been processed, the Web server forgets the client completely.
    - The collection of files that a Web server manages, can be changed without clients having to be informed.

# Stateless vs Stateful servers

- A **stateless server** might have information about clients but no disruption of service if this info is lost.
  - For example, a Web server generally logs all client requests.
  - This information is useful, for example, to decide whether certain documents should be replicated, and where they should be replicated to.
  - Clearly, there is no penalty other than perhaps in the form of suboptimal performance if the log is lost.

# Stateless vs Stateful servers

- A particular form of a stateless design is known as **soft state**.
- In this case, the server promises to maintain state on behalf of the client, but only for a limited time.
- After that time has expired, the server falls back to default behavior, thereby discarding any information it kept on account of the associated client.
- An example of this type of state is:
  - a server promising to keep a client informed about updates, but only for a limited time.
  - After that, the client is required to poll the server for updates.

# Stateless vs Stateful servers

- A **stateful server** generally maintains persistent information on its clients. This means that the information needs to be explicitly deleted by the server.
- e.g., file server:
  - allows a client to keep a local copy of a file
  - maintains a table containing (client, file) entries
  - keeps track of which client currently has the update permissions on which file
  - can improve the performance of read and write operations
  - Drawback: If the server crashes, it has to recover its table of (client, file) entries, or otherwise it cannot guarantee that it has processed the most recent updates on a file.

# Stateless vs Stateful servers

- Session state (temporary):
  - It is associated with a series of operations by a single user and should be maintained for some time, but not indefinitely.
  - It is often maintained in three-tiered client-server architectures, where the application server needs to access a database server through a series of queries before being able to respond to the requesting client.
  - No real harm is done if session state is lost, provided that the client can simply re-issue the original request.
  - Allows for simpler and less reliable storage of state.

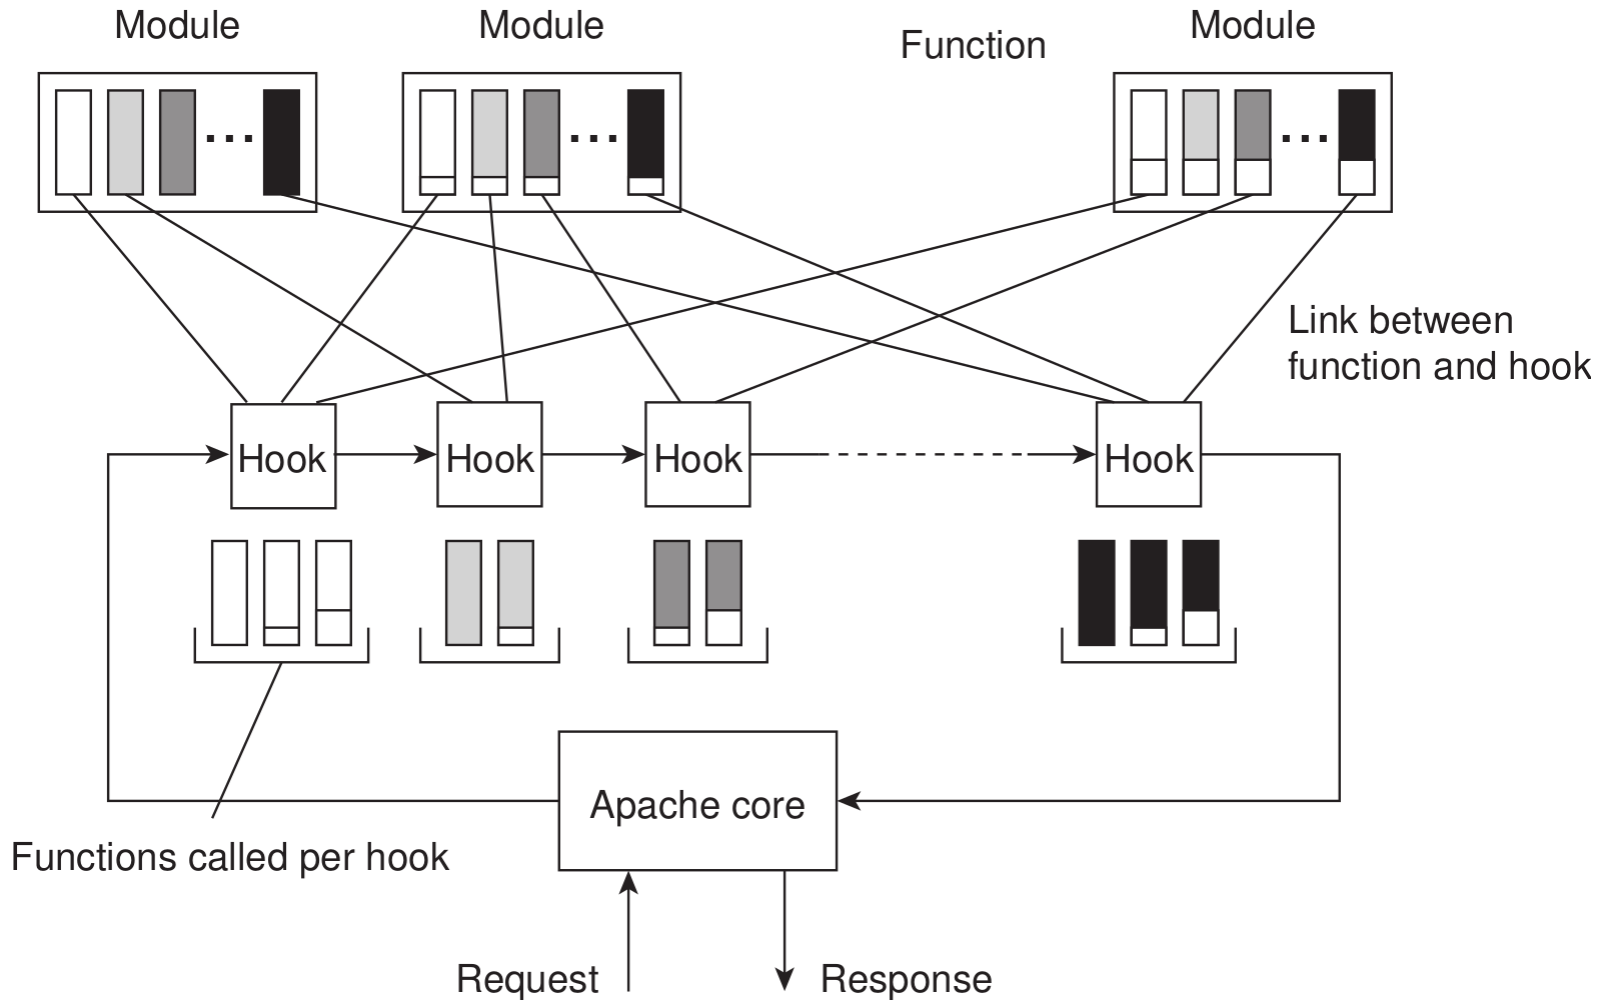
# Stateless vs Stateful servers

- Permanent state:
  - Information maintained in databases e.g., customer information, keys associated with purchased software, etc.
  - Maintaining session state already implies a stateful design
  - Requires special measures when failures do happen
  - Requires explicit assumptions about the durability of state stored at the server.

# Example: The Apache Web server

- Extremely popular server, estimated to be used to host approximately 50% of all Web sites.
- Highly configurable and extensible, and at the same time largely independent of specific platforms.
- Provides its own basic runtime environment, which is then subsequently implemented for different operating systems.
- Apache Portable Runtime (APR):
  - a library that provides a platform-independent interface for file handling, networking, locking, threads, and so on.
  - portability is largely guaranteed provided that only calls to the APR are made and that calls to platform-specific libraries are avoided.

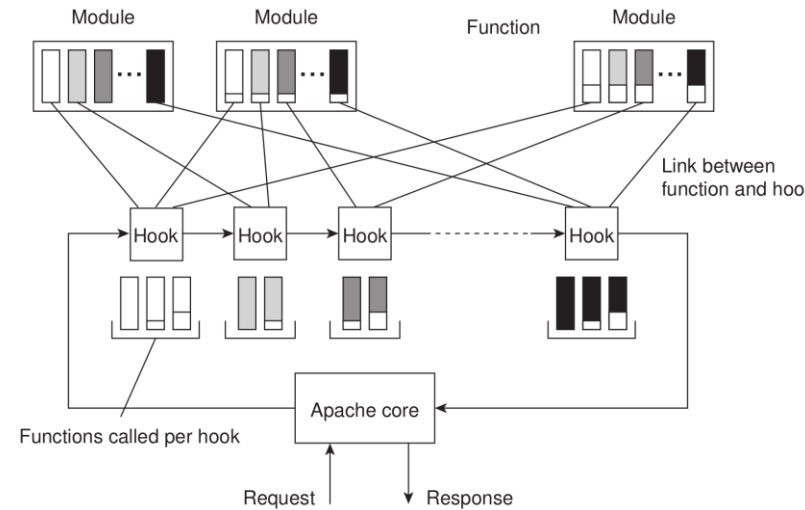
# Example: The Apache Web server





# Example: The Apache Web server

- The functions associated with a hook are all provided by separate **modules**.
- Every hook can contain a set of functions that each should match a specific function prototype (i.e., list of parameters and return type).
- A module developer will write functions for specific hooks.
- When compiling, the developer specifies which function should be added to which hook.



## • Blobs

- Unit of data as key-value pairs
- Value as uninterpreted byte arrays
- Stored internally as a collection of buffers across multiple tiers

## • Bucket

- Collection of blobs
- Flat blob organization

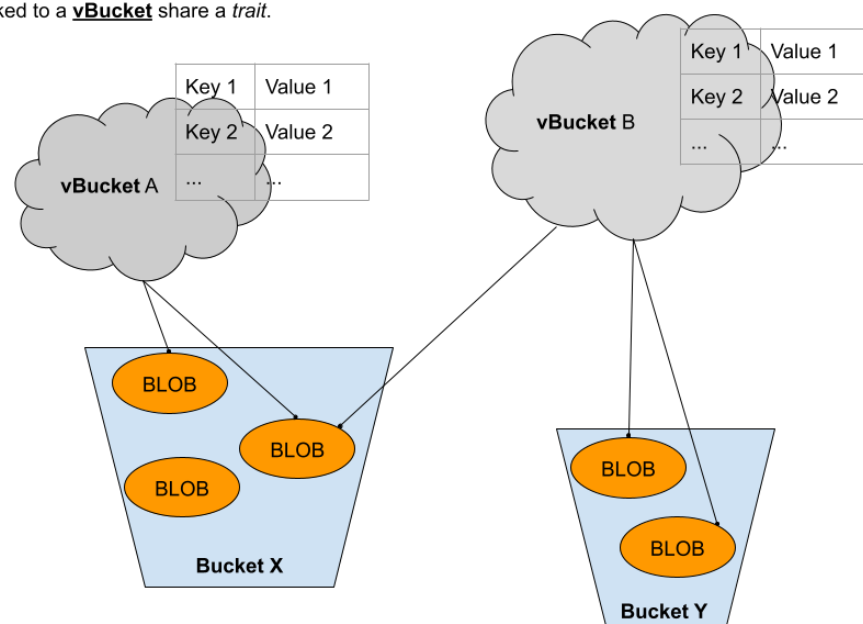
## • Virtual Bucket

- Linked blobs across buckets
- Attached capabilities

## • Traits

- Ordering, grouping, filtering
- Compression, deduplication, etc


Traits represent *capabilities*.  
BLOBs linked to a **vBucket** share a *trait*.



Buckets represent collections of  
(named) **B(L)OBs** (= byte streams).


# Server clusters

- A server cluster is nothing else but a collection of machines connected through a network, where each machine runs one or more servers.



# Request dispatching

- Design goal for server clusters:
  - Hide the fact that there are multiple servers
  - A single access point, implemented through a hardware switch such as a dedicated machine.
  - For scalability and availability, a server cluster may have multiple access points

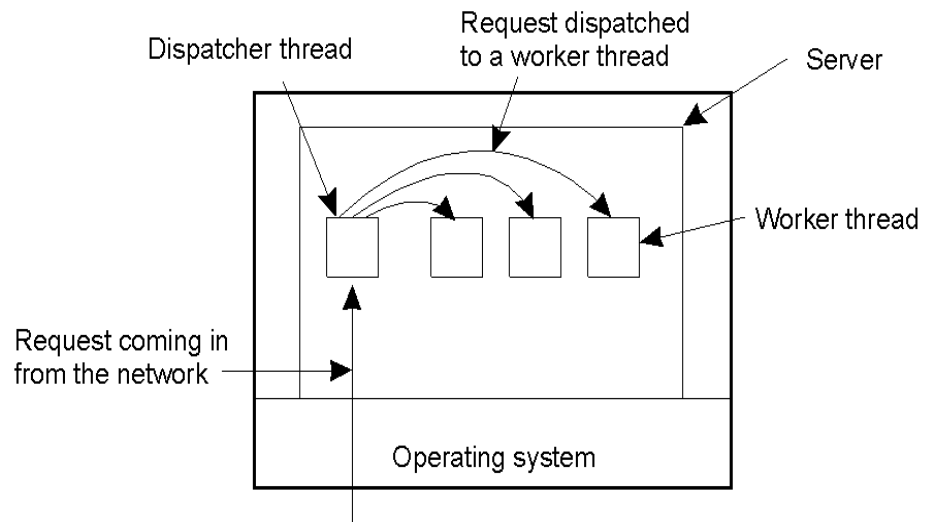


# Multi-threaded Clients Example

- Web Browsers such as IE are multi-threaded
- Such browsers can display data before entire document is downloaded: performs multiple simultaneous tasks
  - Fetch main HTML page, activate separate threads for other parts
  - Each thread sets up a separate connection with the server
    - Uses blocking calls
  - Each part (gif image) fetched separately and in parallel
  - Advantage: connections can be setup to different sources
    - Ad server, image server, web server...

# Multi-threaded Server Example

- Apache web server: pool of pre-spawned worker threads
  - Dispatcher thread waits for requests
  - For each request, choose an idle worker thread
  - Worker thread uses blocking system calls to service web request



# Summary

- Overview of clients and servers (today)
  - Networked user interfaces
  - Thin-client network computing
  - Client-side software for distribution transparency
  - Concurrent versus iterative servers
  - End-points
  - Interrupting a server
  - Stateless vs stateful servers
  - Server clusters
- Readings:
  - Chpt 3.3&3.4 of textbook

# Outline

- Code and process migration
  - Motivation
  - How does migration occur?
  - Resource migration
- Load balancing



# Motivation

- Key reasons: performance and flexibility
- Process migration (aka *strong mobility*)
  - Improved system-wide performance – better utilization of system-wide resources
- Code migration (aka *weak mobility*)
  - Shipment of server code to client
  - Ship parts of client application to server
  - Improve parallelism

# Motivation

- Performance
  - From heavily-loaded to lightly-loaded machines
  - Exploit parallelism, e.g., searching for information in the web through the development of mobile agent, that moves from site to site
  - fault tolerance, e.g., moving from failure-prone to failure-free machines

# Code migration for performance

- Consider a client-server system in which:
  - the server manages a huge database
  - client application needs to perform many database operations involving large quantities of data
  - **better to ship** part of the client application to the server and send only the results across the network
  - Avoid network congestion leading to better performance

# Code migration for performance

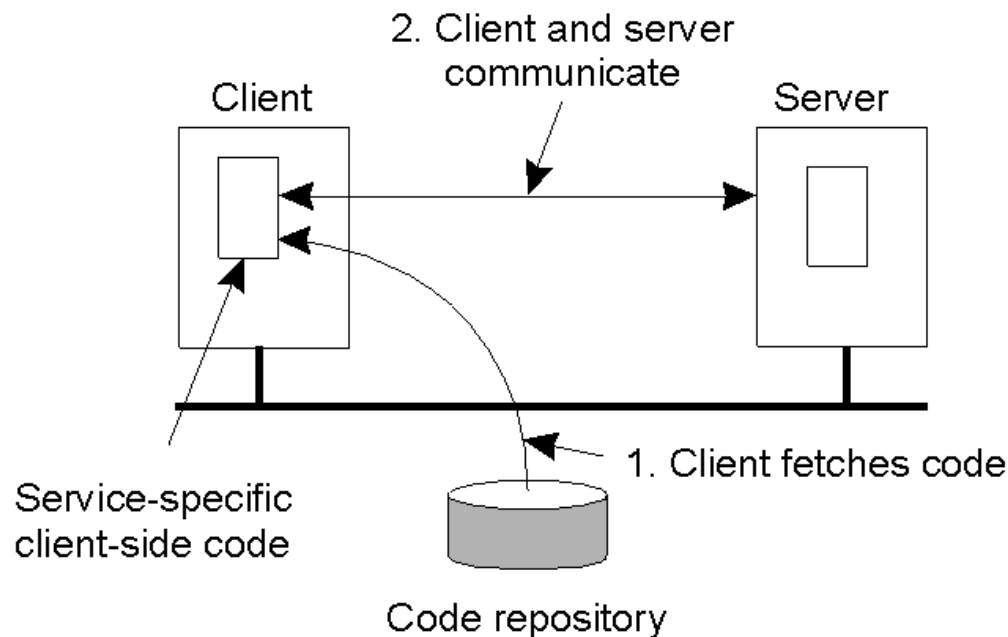
- Consider an interactive database application:
  - clients need to fill in forms that are subsequently translated into a series of database operations
  - Process the form at the client side, and send only the completed form to the server
  - avoid a relatively large number of small messages needing to cross the network
  - client perceives better performance, while at the same time the server spends less time on form processing and communication

# Code migration for performance

- Consider searching the Web:
  - implement a search query in the form of a small mobile program, called **mobile agent**.
  - Agent moves from site to site
  - Linear speedup if we make several copies of such a program, and send each off to different sites
  - Concern: security!

# Motivation

- Flexibility
  - Dynamic configuration of distributed system
  - Clients don't need preinstalled software – download on demand



# Code migration for flexibility

- Suppose a server implements a standardized interface to a file system:
  - server makes use of a proprietary protocol to allow remote clients to access the file system
  - the client-side implementation of the file system interface would need to be linked with the client application
  - Problem: this approach requires that the software be readily available to the client at the time the client application is being developed!

# Code migration for flexibility

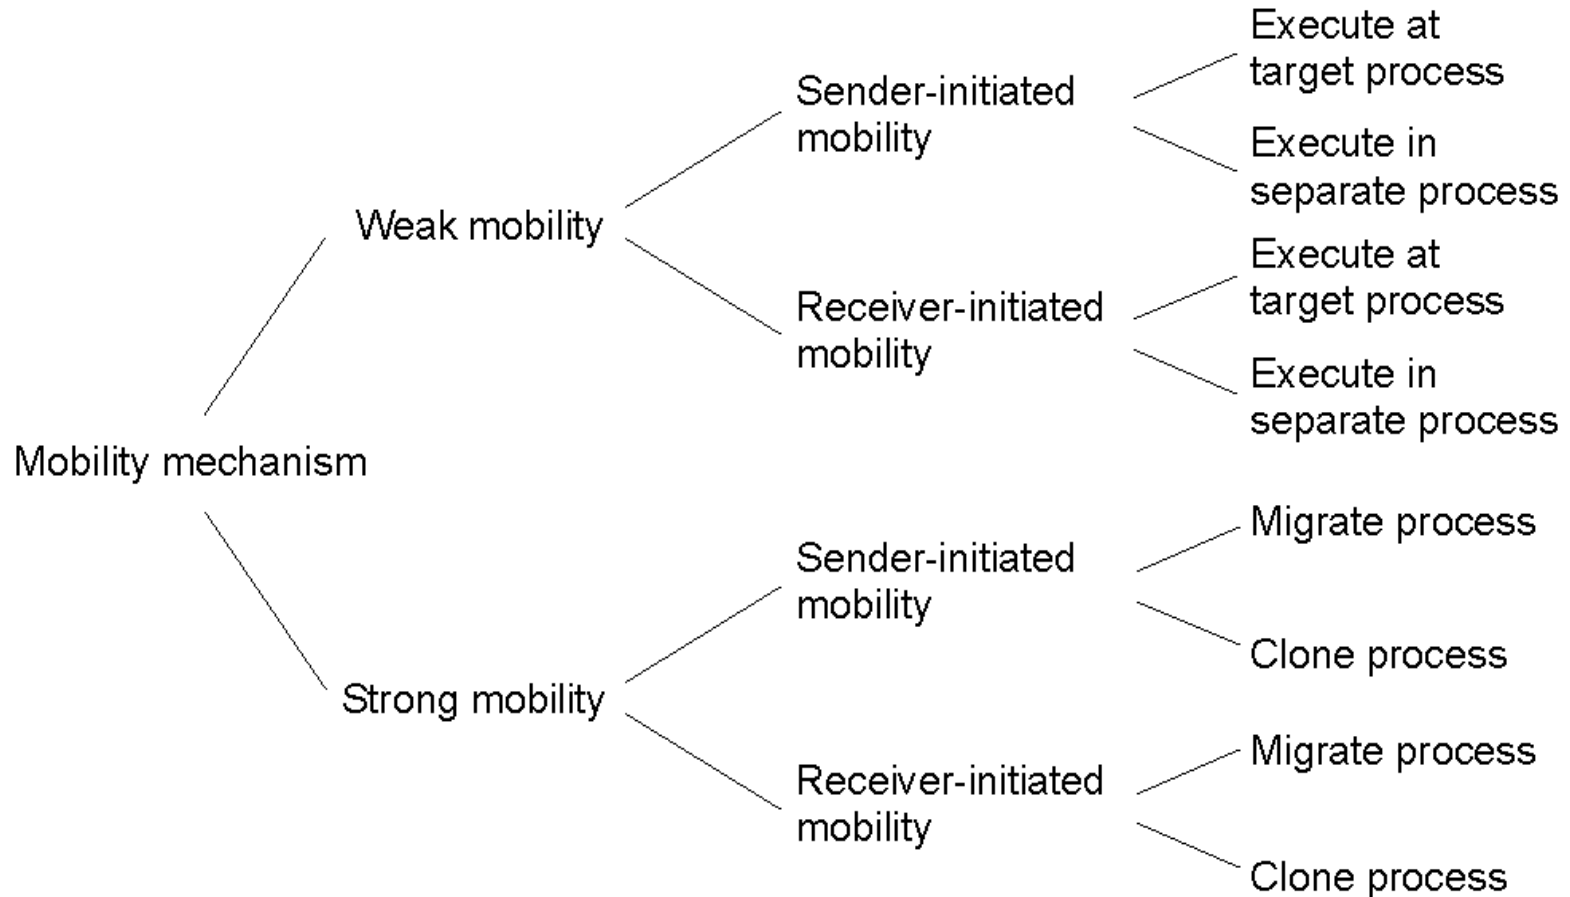
- Alternative:
  - let the server provide the client's implementation when the client binds to the server
  - client dynamically downloads the implementation, goes through the necessary initialization steps, and subsequently invokes the server
  - protocol for downloading and initializing code is standardized
  - necessary that the downloaded code can be executed on the client's machine
  - Benefit: clients need not have all the software preinstalled to talk to servers



# Migration models

- Process = code segment + resource segment + execution segment
- Weak versus strong (runtime info, process) mobility
- Sender-initiated versus receiver-initiated
  - Sender-initiated (code is with sender)
    - Client sending a query to database server
    - Client should be pre-registered
  - Receiver-initiated
    - Java applets
    - Receiver can be anonymous

# Models for Code Migration



# Do Resources Migrate?

- Depends on process-to-resource bindings
  - By identifier: specific web site, ftp server
  - By value: Java libraries
  - By type: printers, local devices
- Depends on resource-to-machine bindings
  - Unattached to any node: data files
  - Fastened resources: database, web sites
  - Fixed resources: local devices, communication end points

# Migration in Heterogeneous Systems

- Systems can be heterogeneous
  - Portability?
- Code migration in heterogeneous systems is being tackled by scripting languages and highly portable languages such as Java
- Rely on a (process) virtual machine that either
  - directly interprets source code (as in the case of scripting languages), or
  - interprets intermediate code generated by a compiler (as in Java).

# Migration in Heterogeneous Systems

- Other solutions:
  - migrate not only processes, but entire computing environments
  - compartmentalize the overall environment and provide processes in the same part their own view on their computing environment
  - Typical example: virtual machines
    - Running an operating system and a suite of applications

# Migration in Heterogeneous Systems

- Migrate VMs:
  - Major advantage is that processes can remain ignorant of the migration itself
  - They need not be interrupted in their execution, nor should they experience any problems with used resources.
  - Resources are either migrating along with a process, or the way that a process accesses a resource is left unaffected.

# Migration in Heterogeneous Systems

- Three ways to handle VM migration:
  - 1) Pushing memory pages to the new machine and resending the ones that are later modified during the migration process.
  - 2) Stopping the current virtual machine; migrate memory, and start the new virtual machine.
  - 3) Letting the new virtual machine pull in new pages as needed, that is, let processes start on the new virtual machine immediately and copy memory pages on demand.